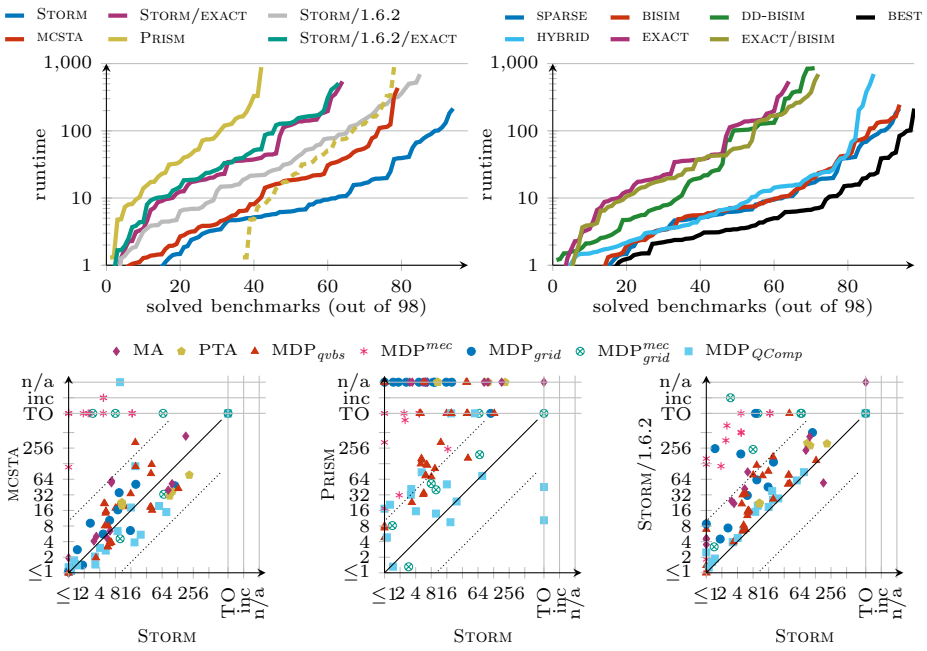


Fig. 6. MDP model checking performance comparison between Storm, mcsta, Prism, and other configurations of Storm. Top: quantile plots indicating how many instances (x-axis) are solved within a given time (y-axis) for each tool (left) and Storm configuration (right). Bottom: scatter plots for direct runtime comparison. n/a means not supported.

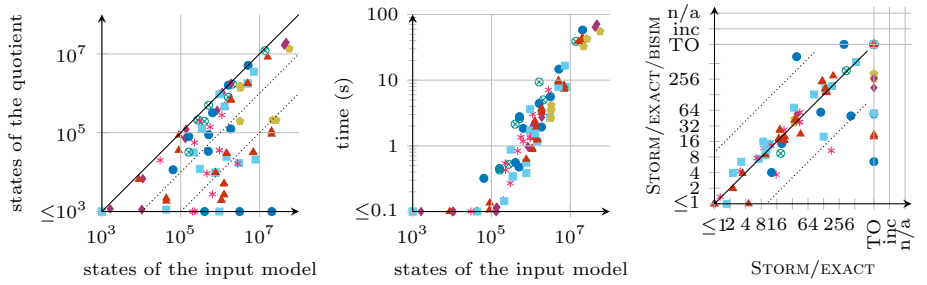


Fig. 7. Bisimulation minimisation in Storm. Left: number of states of the input model (x-axis) and of its bisimulation quotient (y-axis). Middle: time to compute the quotient (y-axis) depending on the number of states of the input model (x-axis). Right: scatter plot for direct runtime comparison of Storm/exact with and without bisimulation.

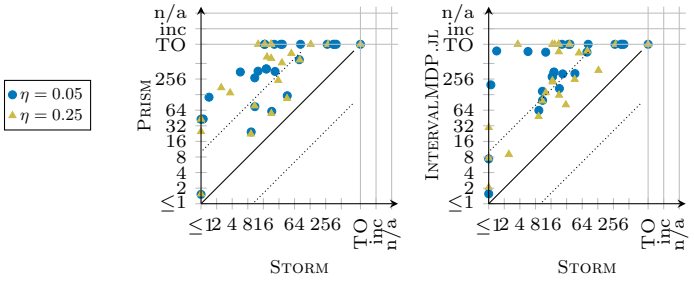


Fig. 8. STORM, PRISM, and INTERVALMDP.JL on interval MDPs with interval width η.

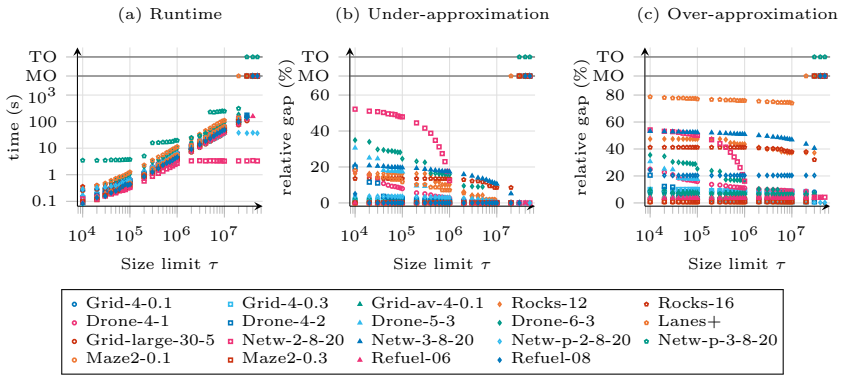


Fig. 9. Belief exploration in STORM with different limits on the number of explored beliefs (size limit τ). Left: runtime. Centre: relative gap between computed bound value and the best computed under-approximation. Right: relative gap between computed bound value and the best computed over-approximation (using belief interpolation). TO and MO indicate timeouts and memory exhaustion.